

03_03_Machine-Level Representation of Programs

why we shouldn't use assembly

- compilers generate pretty fast, efficient code
- tedious, easy to screw up
- not portable

why you shouldn't use assembly

- ... and for that matter, why for a lot of things you shouldn't be using C either ...
- Corbató's Law:
"The number of lines of code a programmer can write in a fixed period of time is the same independent of the language used."

why we're doing assembly here

- learn how the machine works
- what if no compiler?
- processor features not easily accessed by higher level language

the point


- we're not trying to prepare you for a job doing assembly programming
- it's about learning how computers work

how

- look at the assembly generated by GCC
- write some of our own assembly from scratch


```
int sum(int a, int b, int c)
{
    return a+b+c;
}
```

compiler



```
.text
.globl _sum
_sum:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl 12(%ebp), %eax
    addl 8(%ebp), %eax
    addl 16(%ebp), %eax
    leave
    ret
```

assembler



```
0101010110001001111001011000001111101100000010001000101101000101
0000110000000011010001010000100000000011010001010001000011001001
1100001100000000000000000000000000000000000000000000000000000000
0000111100000001000000000000000000000000000000000000000000000000
0000000001011111011100110111010101101101000000000000000000000000
```

instruction set

- set of of operations that a processor supports
- examples
 - load x bytes from this address into register y
 - add what's in register i to what's in register j
- primitive stuff
- usually takes lots of these primitive ops to do something really useful

Other instruction sets

- IA32, Intel64 (x86-64 or x64), IA64
- SPARC
- ARM
- PowerPC
- MIPS
- Alpha
- JVM

- we're using: IA32
 - not easiest, but popular
 - focusing on GCC output on Linux

some terminology

- x86 name for the chips
- IA32 the name of the instruction set
 - IA = Intel **A**rchitecture
- note difference between:
 - **architecture**: what you need to know to program assembly -- instruction set, registers
 - **microarchitecture**: *implementation*
 - *e.g.*, IA32 on non-Intel chips (*e.g.* AMD)

some tools

- compiler – GCC
- assembler – as *aka* gas
- linker – ld
- debugger – gdb
- disassembler - objdump
- profiler - gprof

Some History: Why should we care?

- Important things to take away from the history lesson in the chapter:
 - Moore's law
 - Evolution of register names
 - Backward compatability:
 - Goal: run progs compiled for earlier versions of chip
 - But: old baggage support features that new OS, compilers rarely use


Some History

name	date	transistors	notes
8086	1978	29K	16-bit processor. DOS. 1MB address space. DOS allows 640K
80286	1982	134K	Windows
80386	1985	275K	32-bit registers. Flat addressing. Can run a Unix OS
80486	1989	1.9M	
Pentium	1993	3.1M	
'/MMX	1997	4.5M	instructions helpful for multimedia processing
PentiumPro	1995	6.5M	conditional move instructions
Pentium III	1999	8.2M	
Pentium IV	2001	42M	
Core 2 Duo	2006	291M	
i7	2008	731M	

In parallel. IA64 chips.

name	date	transistors
Itanium	2001	10M
Itanium 2	2002	221M
Itanium 2 Dual-Core	2006	1.7B

Microprocessor Transistor Counts 1971-2011 & Moore's Law



from wikipedia

aside: goals then and now

- big goal then:
 - cram as much processing power on a chip possible
- big goal now:
 - cram as much processing power on a chip possible

BUT

 - don't use so much power
 - some environments: keep the chip small

think cell phone


How bad is your electric bill?

Company	Servers	Electricity	Cost
eBay	16K	$\approx 0.6 \times 10^5$ MWh	$\approx \$3.7\text{M}$
Akamai	40K	$\approx 1.7 \times 10^5$ MWh	$\approx \$10\text{M}$
Rackspace	50K	$\approx 2 \times 10^5$ MWh	$\approx \$12\text{M}$
Microsoft	>200K	$> 6 \times 10^5$ MWh	$> 36\text{M}$
Google	>500K	$> 6.3 \times 10^5$ MWh	$>38\text{M}$
USA (2006)	10.9M	610×10^5 MWh	$\$4.5\text{B}$
MIT campus		2.7×10^5 MWh	$\$62\text{M}$

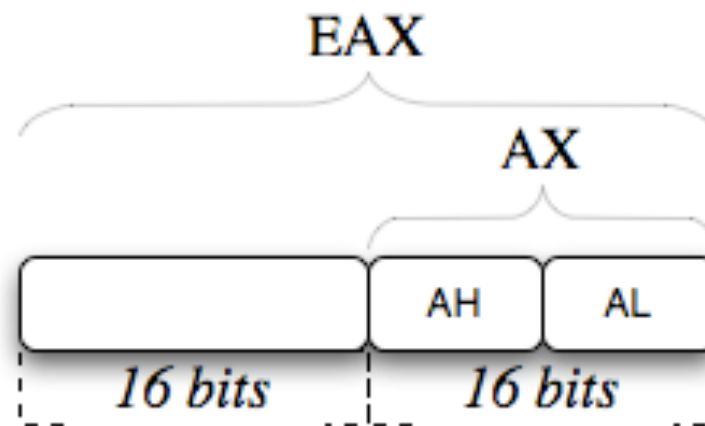
from, *Cutting the Electric Bill for Internet-Scale Systems*, Qureshi et al, CCR 2009.

8086 Register


- Example general purpose register
- 8 general purpose

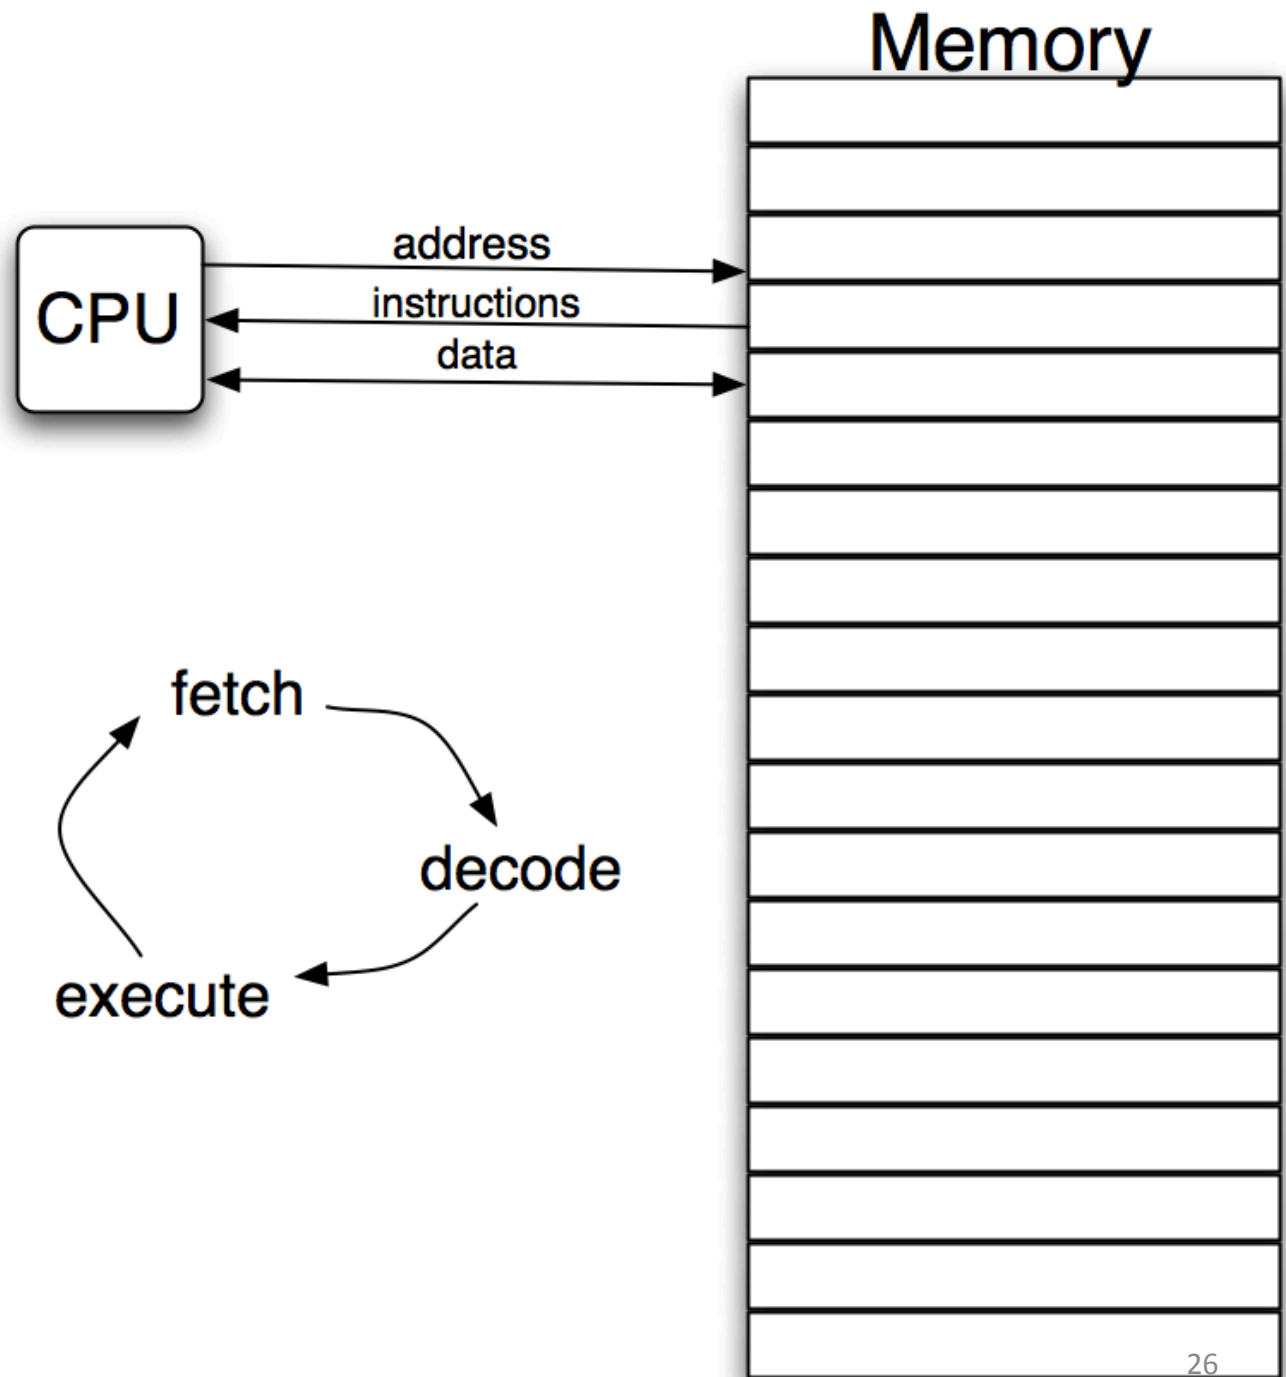


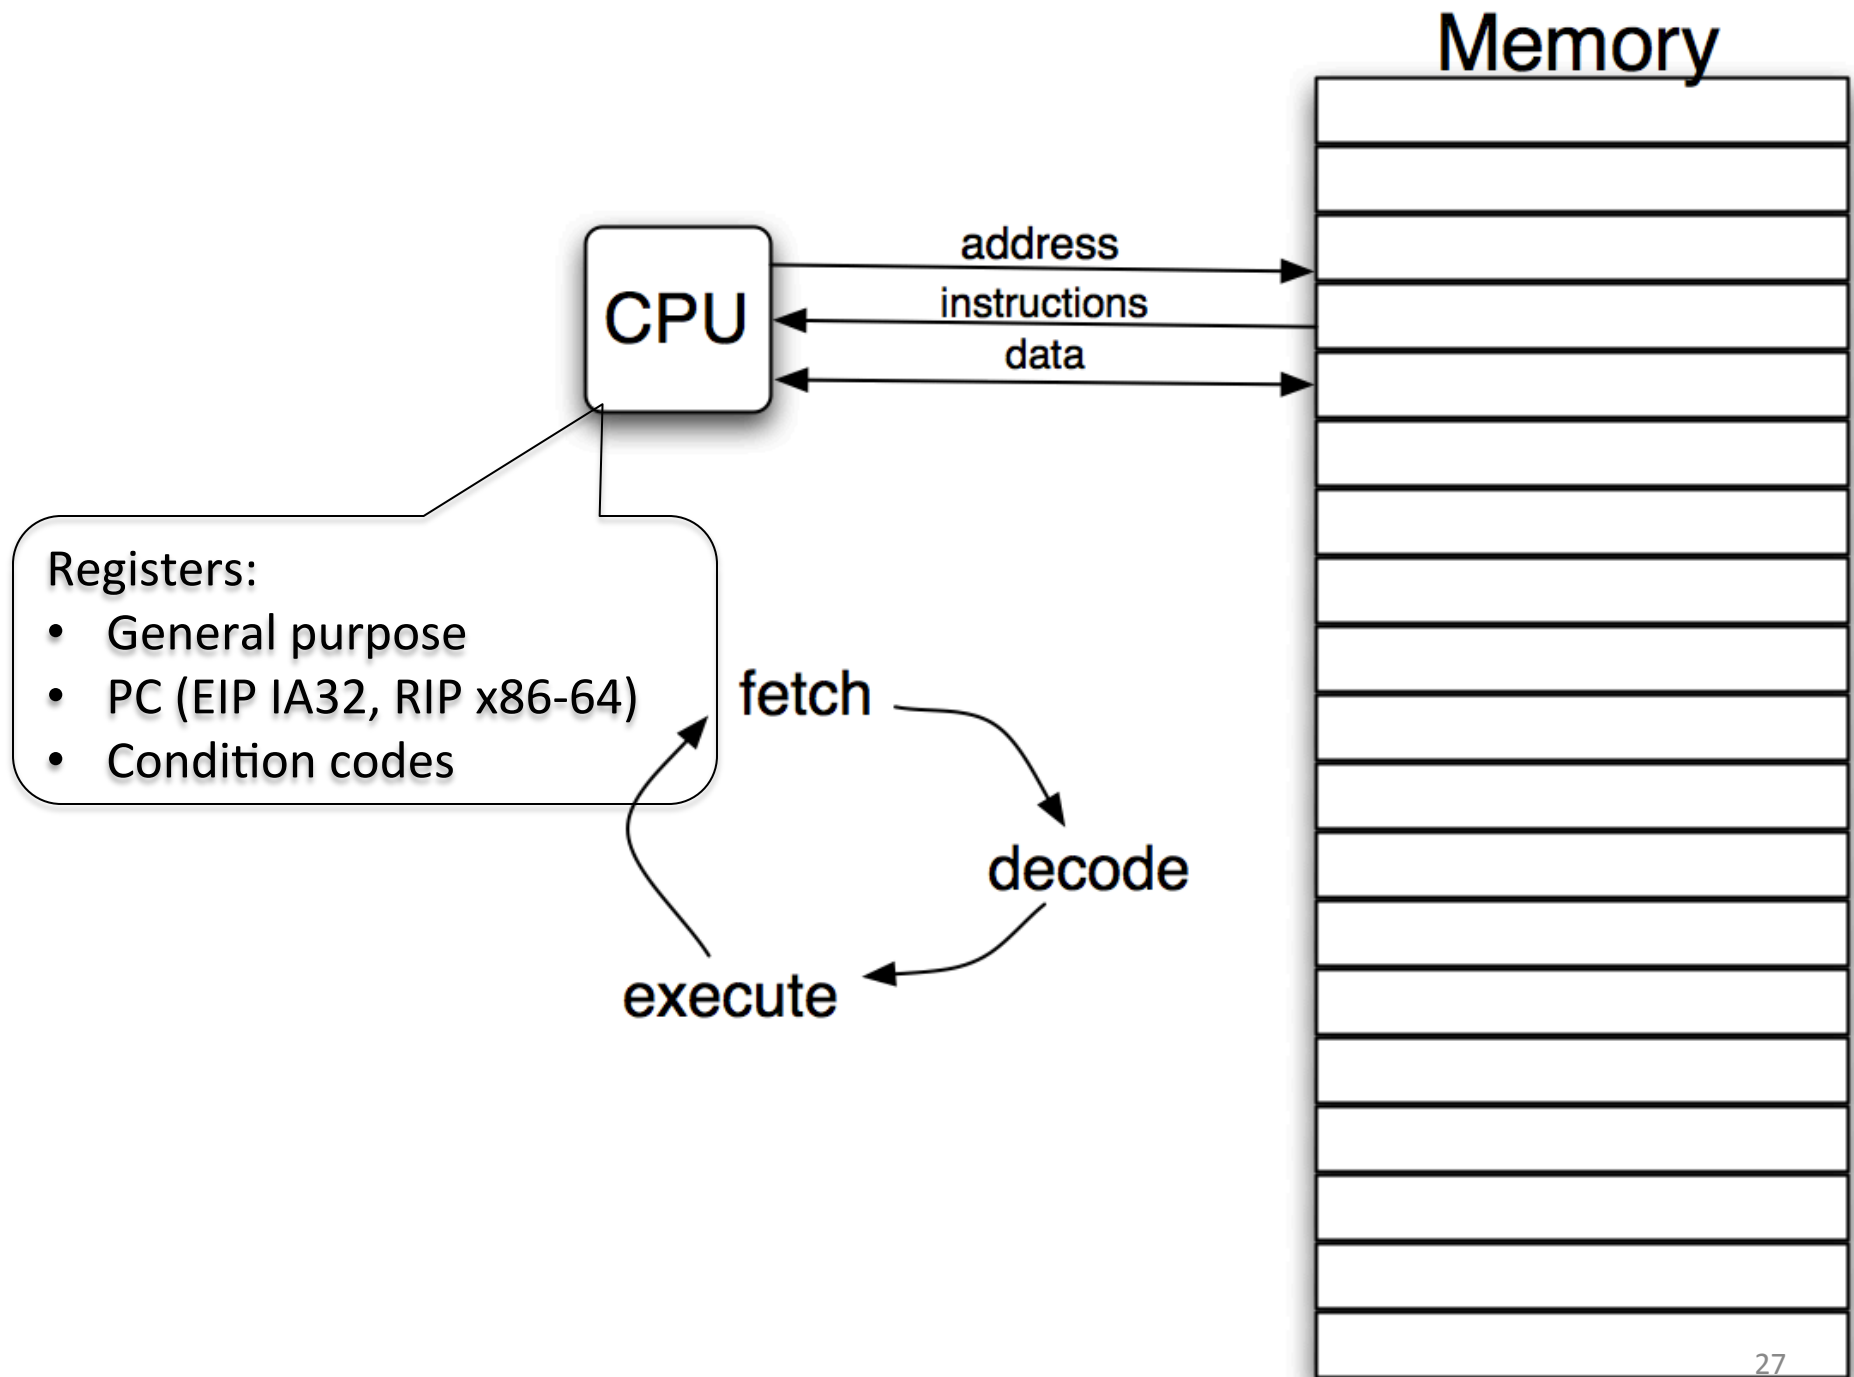
386 registers




64-bit general purpose registers







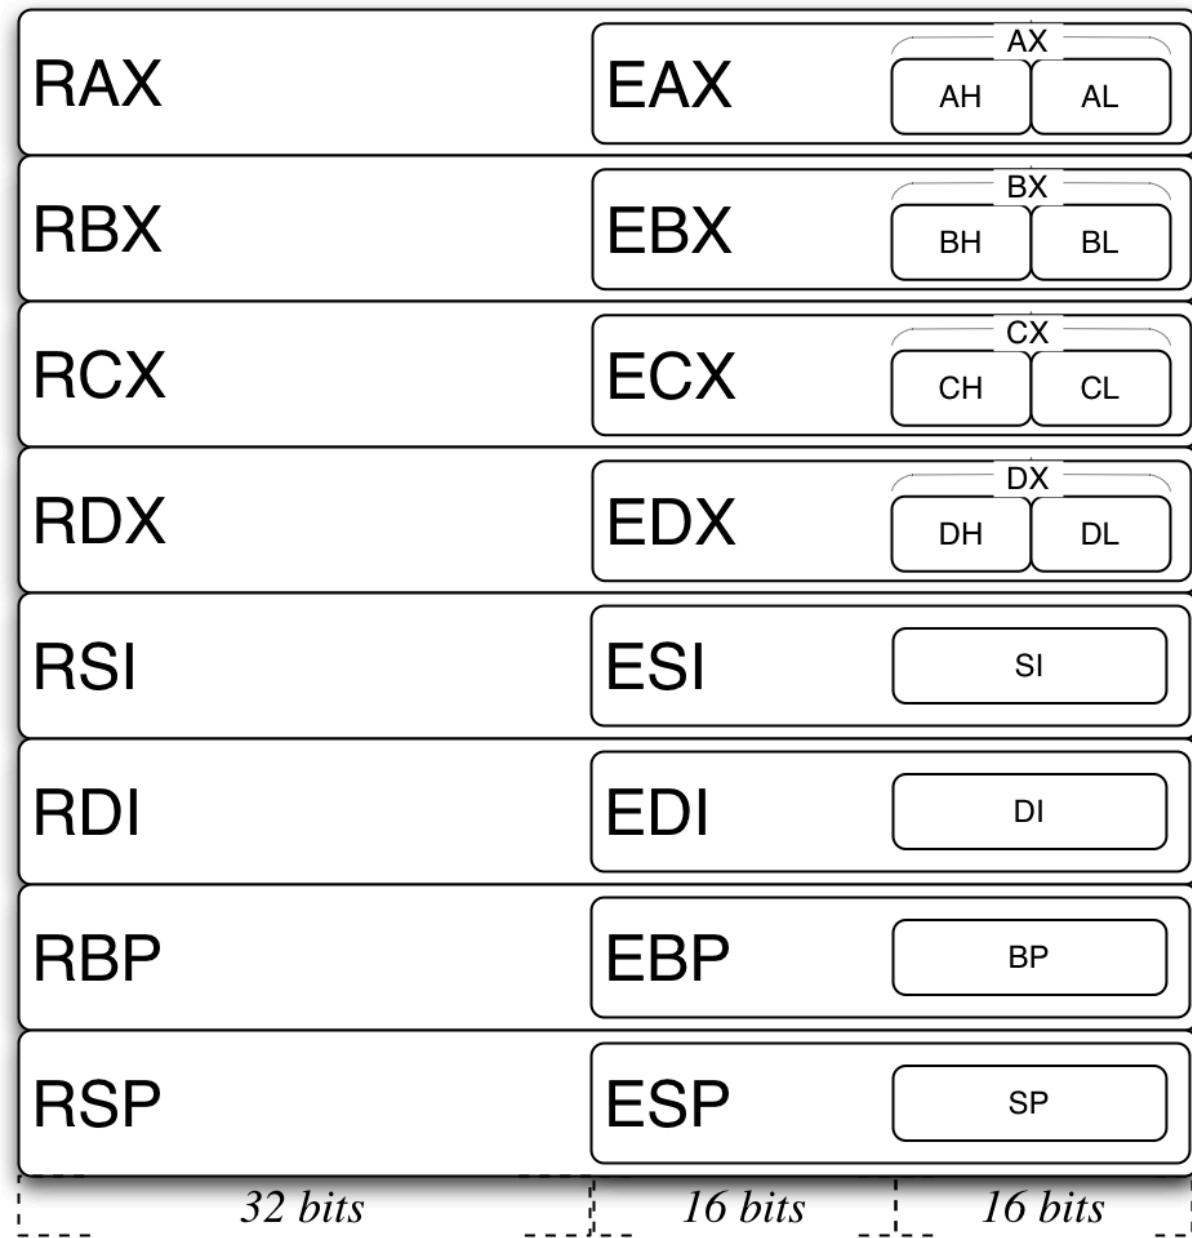
32-bit “Register File”



“general purpose” registers

register	common use
EAX	accumulator, return
EBX	pointer to items in data segment
ECX	loop control
EDX	I/O pointer
ESI	src ptr for string ops
EDI	dst ptr for string ops
ESP	stack pointer
EBP	base pointer

64-bit Register File



but wait, there's more in x86-64

RAX	EAX	<div>AX AH AL</div>	r8	r8d
RBX	EBX	<div>BX BH BL</div>	r9	r9d
RCX	ECX	<div>CX CH CL</div>	r10	r10d
RDX	EDX	<div>DX DH DL</div>	r11	r11d
RSI	ESI	SI	r12	r12d
RDI	EDI	DI	r13	r13d
RBP	EBP	BP	r14	r14d
RSP	ESP	SP	r15	r15d
32 bits			32 bits	32 bits

data types in Intel-speak

name	size
byte	1 byte
word	2 bytes
doubleword	4 bytes
quadword	8 bytes
double quadword	16 bytes

- so a word isn't a word?

simplest function ever

```
1  int sum()  
2  {  
3      int x=30;  
4      int y=57;  
5      int z=39;  
6  
7      return x+y+z;  
8  }
```

gcc output of sum function

```
1          .text
2  .globl _sum
3  _sum:
4          pushl    %ebp                #,
5          movl     %esp, %ebp          #,
6          subl     $24, %esp           #,
7          movl     $30, -20(%ebp)      #, x
8          movl     $57, -16(%ebp)      #, y
9          movl     $39, -12(%ebp)      #, z
10         movl     -16(%ebp), %eax      # y, y
11         addl     -20(%ebp), %eax      # x, D.1509
12         addl     -12(%ebp), %eax      # z, D.1508
13         leave
14         ret
```

simplest function ever

	1	.text
	2	.globl _sum
1	int sum()	3 _sum:
2	{	4 pushl %ebp
3	int x=30;	5 movl %esp, %ebp
4	int y=57;	6 subl \$24, %esp
5	int z=39;	7 movl \$30, -20(%ebp)
6		8 movl \$57, -16(%ebp)
7	return x+y+z;	9 movl \$39, -12(%ebp)
8	}	10 movl -16(%ebp), %eax
		11 addl -20(%ebp), %eax
		12 addl -12(%ebp), %eax
		13 leave
		14 ret

let's try a full program

```
1  /* file summain.c */
2  int main(void)
3  {
4      int x=30;
5      int y=57;
6      int z=39;
7
8      return x+y+z;
9  }
```

its assembly output

```
1  /* file summain.s */
2      .text
3      .globl _main
4  _main:
5      pushl %ebp
6      movl  %esp, %ebp
7      subl  $24, %esp
8      movl  $30, -20(%ebp)
9      movl  $57, -16(%ebp)
10     movl  $39, -12(%ebp)
11     movl  -16(%ebp), %eax
12     addl  -20(%ebp), %eax
13     addl  -12(%ebp), %eax
14     leave
15     ret
16     .subsections_via_symbols
```


same code
on SPARC

```
1  .section          ".text"
2      .align 4
3      .global main
4      .type      main,#function
5      .proc      04
6  main:
7      !#PROLOGUE# 0
8      save      %sp, -128, %sp
9      !#PROLOGUE# 1
10     mov      30, %o0
11     st       %o0, [%fp-20]
12     mov      57, %o0
13     st       %o0, [%fp-24]
14     mov      39, %o0
15     st       %o0, [%fp-28]
16     ld       [%fp-20], %o0
17     ld       [%fp-24], %o1
18     add      %o0, %o1, %o0
19     ld       [%fp-28], %o1
20     add      %o0, %o1, %o0
21     mov      %o0, %i0
22     b        .LL2
23     nop
24 .LL2:
25     ret
26     restore
27 .LLfe1:
28     .size      main,.LLfe1-main
29     .ident      "GCC: (GNU) 2.95.3 20010315 (release)"
```

*most of
the same
on ARM*

```
1  sum:
2      @ args = 0, pretend = 0, frame = 16
3      @ frame_needed = 1, uses_anonymous_args = 0
4      @ link register save eliminated.
5      push        {r7}
6      sub         sp, sp, #20
7      add         r7, sp, #0
8      mov         r3, #30
9      str         r3, [r7, #12]
10     mov         r3, #57
11     str         r3, [r7, #8]
12     mov         r3, #39
13     str         r3, [r7, #4]
14     ldr         r2, [r7, #12]
15     ldr         r3, [r7, #8]
16     adds        r2, r2, r3
17     ldr         r3, [r7, #4]
18     adds        r3, r2, r3
19     mov         r0, r3
20     add         r7, r7, #20
21     mov         sp, r7
22     pop         {r7}
23     bx          lr
```

sum.C

```
int sum(int a, int b, int c)
{
    return a+b+c;
}
```

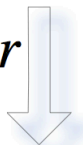
compiler



sum.S

```
.text
.globl _sum
_sum:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl 12(%ebp), %eax
    addl 8(%ebp), %eax
    addl 16(%ebp), %eax
    leave
    ret
```

assembler



sum.o

```
0101010110001001111001011000001111101100000010001000101101000101
0000110000000011010001010000100000000011010001010001000011001001
110000110000000000000000000000000000000000000000000000000000000
000011110000000100000000000000000000000000000000000000000000000
0000000001011111011100110111010101101101000000000000000000000000
```

executable file:

sum

```
0101010110001001111001011000001111101100000010001000101101000101
0000110000000011010001010000100000000011010001010001000011001001
110000110000000000000000000000000000000000000000000000000000000
000011110000000100000000000000000000000000000000000000000000000
0000000001011111011100110111010101101101000000000000000000000000
```

linker




libc

```
printf
scanf
fopen
...
strcpy
...
```

sum.C

```
int sum(int a, int b, int c)
{
    return a+b+c;
}
```

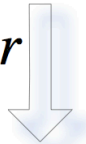
compiler



sum.S

```
.text
.globl _sum
_sum:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl 12(%ebp), %eax
    addl 8(%ebp), %eax
    addl 16(%ebp), %eax
    leave
    ret
```

assembler



sum.o

```
0101010110001001111001011000001111101100000010001000101101000101
0000110000000011010001010000100000000011010001010001000011001001
1100001100000000000000000000000000000000000000000000000000000000
0000111100000001000000000000000000000000000000000000000000000000
0000000001011111011100110111010101101101000000000000000000000000
```

executable file:

sum

```
0101010110001001111001011000001111101100000010001000101101000101
0000110000000011010001010000100000000011010001010001000011001001
1100001100000000000000000000000000000000000000000000000000000000
0000111100000001000000000000000000000000000000000000000000000000
0000000001011111011100110111010101101101000000000000000000000000
```

understand what's in each

libc

```
printf
scanf
fopen
...
strcpy
...
```

linker



sum.C

```
int sum(int a, int b, int c)
{
    return a+b+c;
}
```

decompiler?



opposite direction?

sum.S

```
.text
.globl _sum
_sum:
    pushl %ebp
    movl %esp, %ebp
    subl $8, %esp
    movl 12(%ebp), %eax
    addl 8(%ebp), %eax
    addl 16(%ebp), %eax
    leave
    ret
```

disassembler



sum.O

```
0101010110001001111001011000001111101100000010001000101101000101
0000110000000011010001010000100000000011010001010001000011001001
1100001100000000000000000000000000000000000000000000000000000000
0000111100000001000000000000000000000000000000000000000000000000
0000000001011110111001101110101011011010000000000000000000000000
```

what can we disassemble?

- any executable
- disassembler interprets bytes as assembly src
- no source code required

binary compatibility

- Linux runs on my Intel desktop
- Windows runs on my Intel desktop
- Why can't I take my Windows binaries and run them on Linux and vice versa?